

# Ch 5 1 ModelSelect

Analyze the ISLR (Introduction to Statistical Learning with R) data package's baseball 'Hitters' data frame:

```
library(ISLR)
summary(Hitters)
```

```
##           AtBat           Hits           HmRun           Runs
## Min.      : 16.0   Min.       :  1   Min.       : 0.00   Min.       :  0.00
## 1st Qu.:255.2   1st Qu.: 64   1st Qu.: 4.00   1st Qu.: 30.25
## Median :379.5   Median : 96   Median : 8.00   Median : 48.00
## Mean     :380.9   Mean    :101   Mean     :10.77   Mean     : 50.91
## 3rd Qu.:512.0   3rd Qu.:137   3rd Qu.:16.00   3rd Qu.: 69.00
## Max.     :687.0   Max.     :238   Max.     :40.00   Max.     :130.00
##
##           RBI           Walks           Years           CAtBat
## Min.      :  0.00   Min.       :  0.00   Min.       : 1.000   Min.       :  19.0
## 1st Qu.: 28.00   1st Qu.: 22.00   1st Qu.: 4.000   1st Qu.: 816.8
## Median : 44.00   Median : 35.00   Median : 6.000   Median :1928.0
## Mean     : 48.03   Mean     : 38.74   Mean     : 7.444   Mean     :2648.7
## 3rd Qu.: 64.75   3rd Qu.: 53.00   3rd Qu.:11.000   3rd Qu.:3924.2
## Max.     :121.00   Max.      :105.00   Max.      :24.000   Max.     :14053.0
##
##           CHits           CHmRun           CRuns           CRBI
## Min.      :  4.0   Min.       :  0.00   Min.       :  1.0   Min.       :  0.00
## 1st Qu.: 209.0   1st Qu.: 14.00   1st Qu.: 100.2   1st Qu.: 88.75
## Median : 508.0   Median : 37.50   Median : 247.0   Median :220.50
## Mean     : 717.6   Mean     : 69.49   Mean     : 358.8   Mean     :330.12
## 3rd Qu.:1059.2   3rd Qu.: 90.00   3rd Qu.: 526.2   3rd Qu.:426.25
## Max.     :4256.0   Max.      :548.00   Max.      :2165.0   Max.     :1659.00
##
##           CWalks           League Division           PutOuts           Assists
## Min.      :  0.00   A:175   E:157   Min.       :  0.0   Min.       :  0.0
## 1st Qu.: 67.25   N:147   W:165   1st Qu.: 109.2   1st Qu.:  7.0
## Median : 170.50                               Median : 212.0   Median : 39.5
## Mean     : 260.24                               Mean     : 288.9   Mean     :106.9
## 3rd Qu.: 339.25                               3rd Qu.: 325.0   3rd Qu.:166.0
## Max.     :1566.00                               Max.      :1378.0   Max.     :492.0
##
##           Errors           Salary           NewLeague
## Min.      :  0.00   Min.       : 67.5   A:176
## 1st Qu.:  3.00   1st Qu.: 190.0   N:146
## Median :  6.00   Median : 425.0
## Mean     :  8.04   Mean     :535.9
## 3rd Qu.: 11.00   3rd Qu.: 750.0
## Max.     : 32.00   Max.     :2460.0
##                      NA's      :59
```

There are missing values, before we proceed we will remove them:

```
with(Hitters, sum(is.na(Salary)))
```

```
## [1] 59
```

```
Hitters=na.omit(Hitters)
with(Hitters, sum(is.na(Salary)))
```

```
## [1] 0
```

## Best Subset regression

We will now use the package `leaps` to evaluate all the best-subset models.

```
library(leaps)
regfit.full = regsubsets(Salary~., data=Hitters)
summary(regfit.full)
```

```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., data = Hitters)
## 19 Variables (and intercept)
##           Forced in Forced out
## AtBat      FALSE      FALSE
## Hits       FALSE      FALSE
## HmRun      FALSE      FALSE
## Runs       FALSE      FALSE
## RBI        FALSE      FALSE
## Walks      FALSE      FALSE
## Years      FALSE      FALSE
## CAtBat     FALSE      FALSE
## CHits      FALSE      FALSE
## CHmRun     FALSE      FALSE
## CRuns      FALSE      FALSE
## CRBI       FALSE      FALSE
## CWalks     FALSE      FALSE
## LeagueN    FALSE      FALSE
## DivisionW  FALSE      FALSE
## PutOuts    FALSE      FALSE
## Assists    FALSE      FALSE
## Errors     FALSE      FALSE
## NewLeagueN FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun CRuns
## 1  ( 1 ) " "   " "   " "   " "   " "   " "   " "   " "   " "
## 2  ( 1 ) " "   "*"  " "   " "   " "   " "   " "   " "   " "
## 3  ( 1 ) " "   "*"  " "   " "   " "   " "   " "   " "   " "
## 4  ( 1 ) " "   "*"  " "   " "   " "   " "   " "   " "   " "
## 5  ( 1 ) "*"   "*"  " "   " "   " "   " "   " "   " "   " "
## 6  ( 1 ) "*"   "*"  " "   " "   " "   "*"  " "   " "   " "   " "
## 7  ( 1 ) " "   "*"  " "   " "   " "   "*"  " "   "*"   "*"   " "
## 8  ( 1 ) "*"   "*"  " "   " "   " "   "*"  " "   " "   "*"   "*"

```

```
##          CRBI CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) "*" " " " " " " " " " " " "
## 2 ( 1 ) "*" " " " " " " " " " " " "
## 3 ( 1 ) "*" " " " " " " "*" " " " " "
## 4 ( 1 ) "*" " " " " "*" "*" " " " " " "
## 5 ( 1 ) "*" " " " " "*" "*" " " " " " "
## 6 ( 1 ) "*" " " " " "*" "*" " " " " " "
## 7 ( 1 ) " " " " " " "*" "*" " " " " " "
## 8 ( 1 ) " " "*" " " "*" "*" " " " " " "
```

By default, it gives the first 8 variables best-subset models. Let's do it again for all the variables:

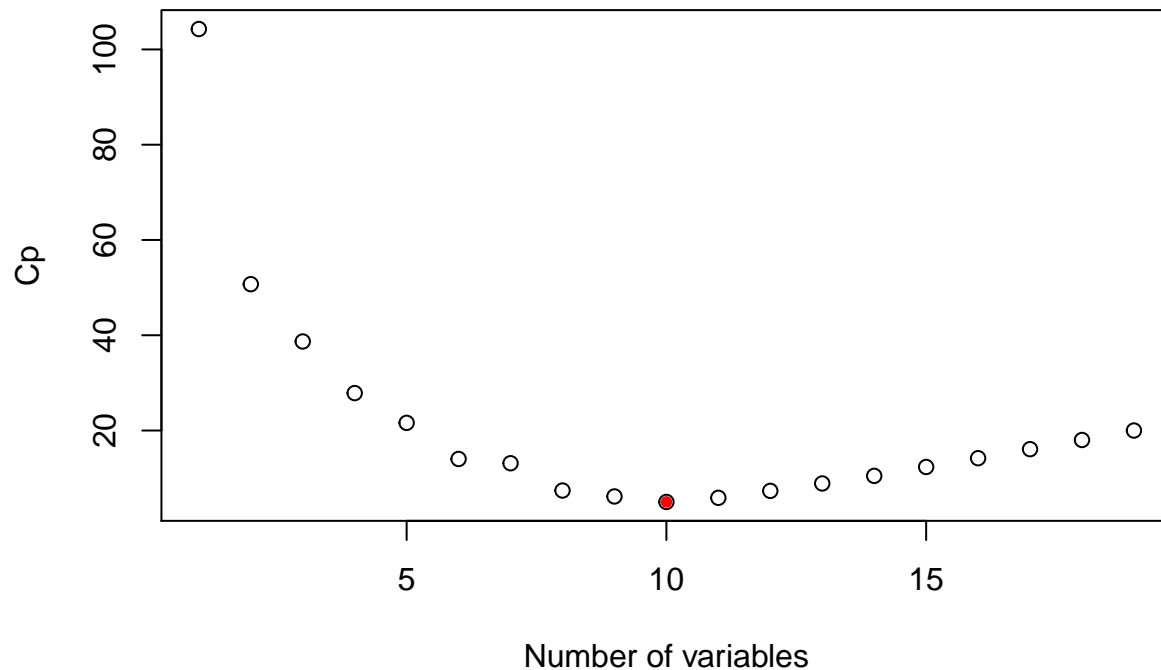
```
regfit.full = regsubsets(Salary~., data=Hitters, nvmax=19)
reg.summary = summary(regfit.full)
names(reg.summary)
```

```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
plot(reg.summary$cp, xlab="Number of variables", ylab="Cp")
which.min(reg.summary$cp)
```

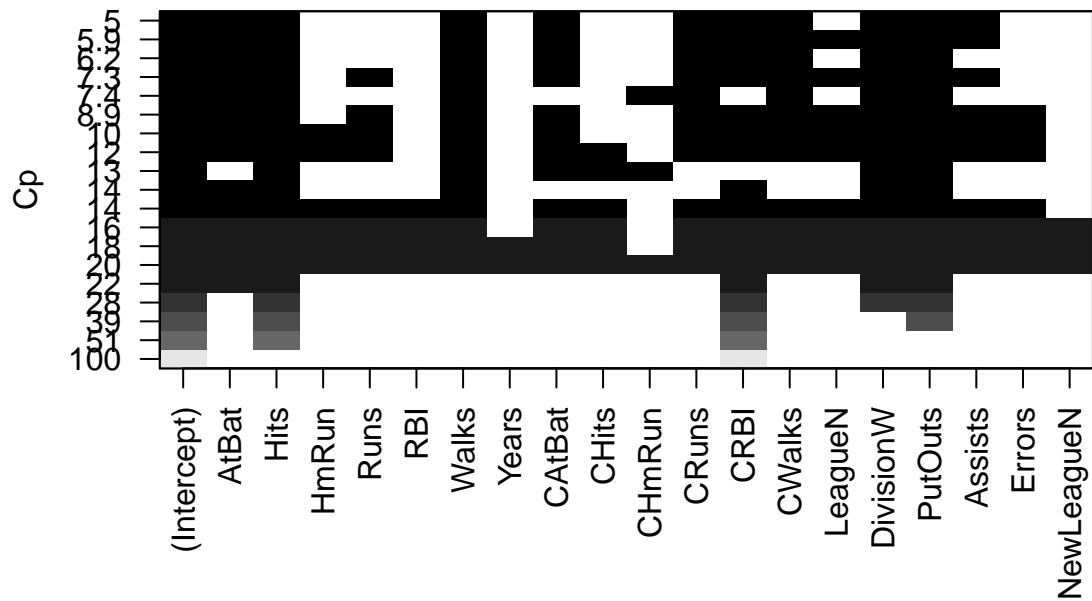
```
## [1] 10
```

```
points(10, reg.summary$cp[10], pch=20, col="red")
```



There is a method for the regsubset object:

```
plot(regfit.full, scale="Cp")
```



```
coef(regfit.full, 10)
```

```
## (Intercept)      AtBat      Hits      Walks      CAtBat
## 162.5354420    -2.1686501    6.9180175    5.7732246   -0.1300798
##      CRuns      CRBI      CWalks    DivisionW    PutOuts
##   1.4082490    0.7743122   -0.8308264  -112.3800575    0.2973726
##      Assists
##   0.2831680
```

## Forward Stepwise Selection

We use `regsubset` again, but specify the `method="forward"` option.

```
regfit.fwd = regsubsets(Salary~., data=Hitters, nvmax=19, method="forward")
summary(regfit.fwd)
```

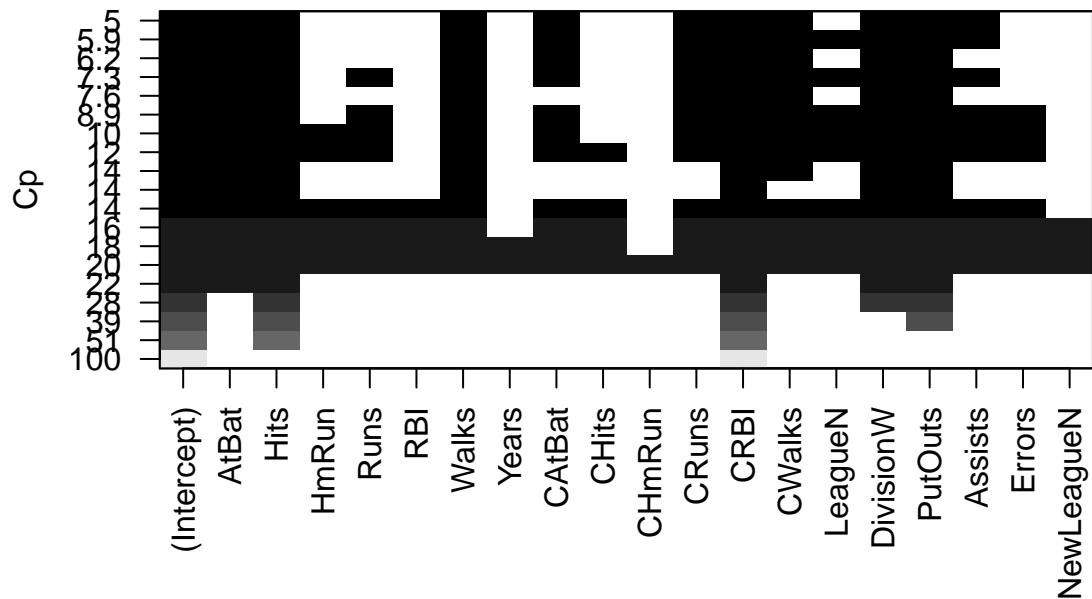
```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., data = Hitters, nvmax = 19, method = "forward")
## 19 Variables (and intercept)
##           Forced in Forced out
## AtBat      FALSE      FALSE
## Hits       FALSE      FALSE
## HmRun      FALSE      FALSE
## Runs       FALSE      FALSE
## RBI        FALSE      FALSE
## Walks      FALSE      FALSE
## Years      FALSE      FALSE
## CAtBat     FALSE      FALSE
## CHits      FALSE      FALSE
## CHmRun     FALSE      FALSE
## CRuns      FALSE      FALSE
## CRBI       FALSE      FALSE
```

```

## CWalks          FALSE      FALSE
## LeagueN         FALSE      FALSE
## DivisionW       FALSE      FALSE
## PutOuts         FALSE      FALSE
## Assists         FALSE      FALSE
## Errors          FALSE      FALSE
## NewLeagueN      FALSE      FALSE
## 1 subsets of each size up to 19
## Selection Algorithm: forward
##      AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun CRuns
## 1 ( 1 ) " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " "*" " " " " " " " " " " " " " "
## 3 ( 1 ) " " "*" " " " " " " " " " " " " " "
## 4 ( 1 ) " " "*" " " " " " " " " " " " " " "
## 5 ( 1 ) "*" "*" " " " " " " " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "
## 7 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "
## 8 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*"
## 9 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " "*"
## 10 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " "*"
## 11 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " "*"
## 12 ( 1 ) "*" "*" " " " " "*" " " "*" " " " "*"
## 13 ( 1 ) "*" "*" " " " " "*" " " "*" " " " "*"
## 14 ( 1 ) "*" "*" "*" "*" " " "*" " " "*" " " " "*"
## 15 ( 1 ) "*" "*" "*" "*" " " "*" " " "*" "*" " " "*"
## 16 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" "*" " " "*"
## 17 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" "*" " " "*"
## 18 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" " " "*"
## 19 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
##      CRBI CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) "*" " " " " " " " " " " " "
## 2 ( 1 ) "*" " " " " " " " " " " " "
## 3 ( 1 ) "*" " " " " " " "*" " " " " "
## 4 ( 1 ) "*" " " " " "*" "*" " " " " "
## 5 ( 1 ) "*" " " " " "*" "*" " " " " "
## 6 ( 1 ) "*" " " " " "*" "*" " " " " "
## 7 ( 1 ) "*" "*" " " "*" "*" " " " " "
## 8 ( 1 ) "*" "*" " " "*" "*" " " " " "
## 9 ( 1 ) "*" "*" " " "*" "*" " " " " "
## 10 ( 1 ) "*" "*" " " "*" "*" "*" " " " "
## 11 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "
## 12 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
## 14 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
## 15 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
## 16 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
## 17 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"
## 18 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"
## 19 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"

```

```
plot(regfit.fwd, scale="Cp")
```



## Model Selection Using a Validation Set

Let's make a training and validation set, so that we can choose a good subset model.

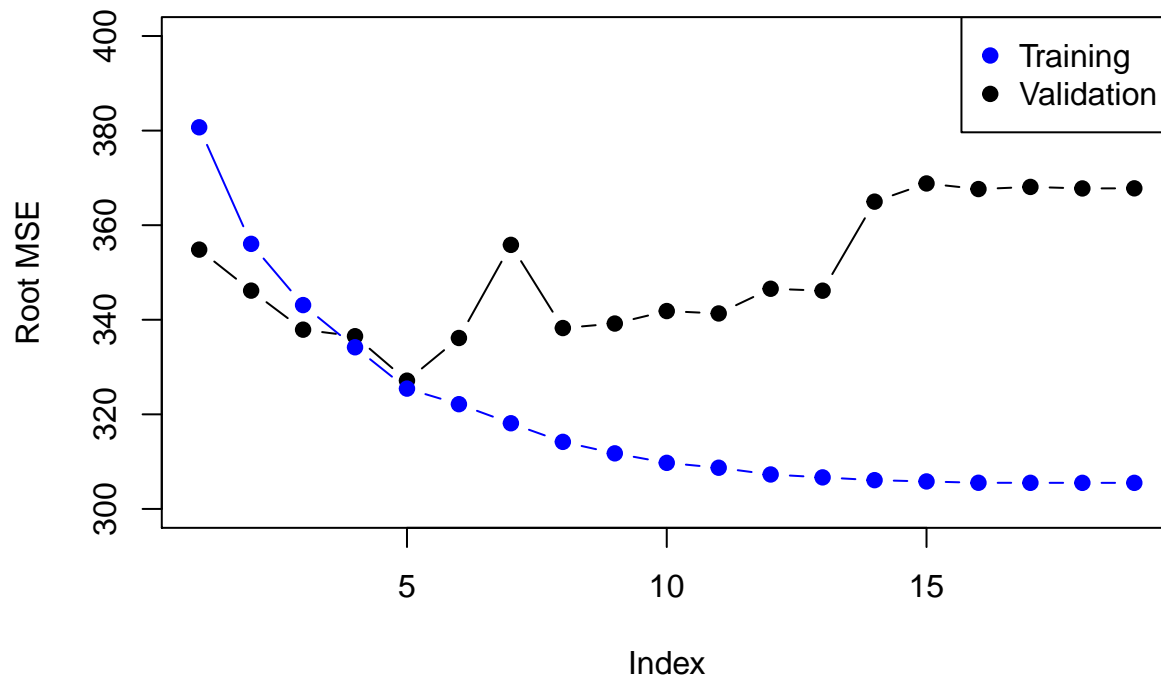
```
dim(Hitters)
```

```
## [1] 263 20
```

```
set.seed(1)
train = sample(seq(263),180,replace=FALSE)
regfit.fwd = regsubsets(Salary~., data=Hitters[train,], nvmax=19, method="forward")
```

Now, we separate the data to two parts, one for training and another for test/validation to make prediction. There is no `prediction` method for `regsubsets`, so we need to write that part. We also create a vector to store the results of the 19 different models.

```
val.errors = rep(NA, 19)
x.test = model.matrix(Salary~., data=Hitters[-train,])
for(i in 1:19){
  coefi = coef(regfit.fwd, id=i)
  pred = x.test[,names(coefi)]%*%coefi
  val.errors[i] = mean((Hitters$Salary[-train]-pred)^2)
}
plot(sqrt(val.errors), ylab="Root MSE", ylim=c(300,400), pch=19, type="b")
points(sqrt(regfit.fwd$rss[-1]/180),col="blue",pch=19,type="b") # -1 excludes null model
legend("topright", legend=c("Training","Validation"),col=c("blue","black"),pch=19)
```



As expected, the model error goes down monotonically as the model gets bigger, but not so for the validation error.

This was a little tedious - not having a `predict` method `regsubsets`. So we will write one!

```
predict.regsubsets = function(object, newdata, id, ...){
  form = as.formula(object$call[[2]])
  mat = model.matrix(form, newdata)
  coefi = coef(object, id=id)
  mat[, names(coefi)] %*% coefi
}
as.formula(regfit.fwd$call[[2]])
```

```
## Salary ~ .
```

## Model Seletion by Cross-Validation

We will do a 10-fold cross-validation.

```
set.seed(11)
folds = sample(rep(1:10,length=nrow(Hitters)))
table(folds)
```

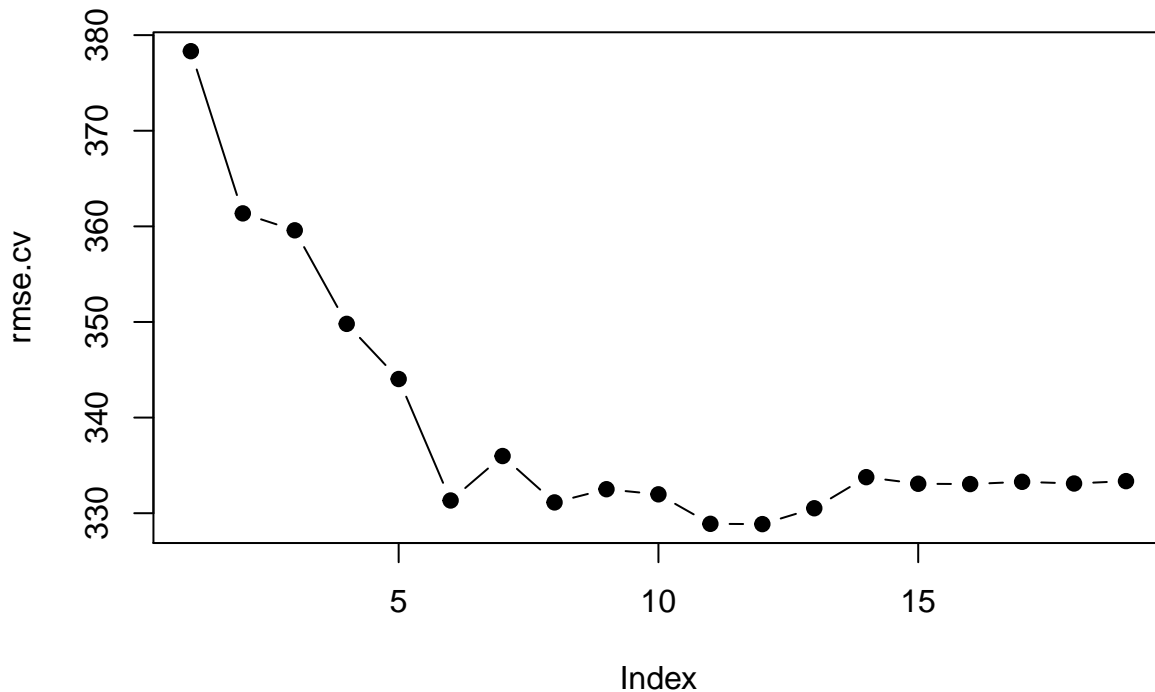
```
## folds
##  1  2  3  4  5  6  7  8  9 10
## 27 27 27 26 26 26 26 26 26 26
```

```
cv.errors = matrix(NA,10,19)
for(k in 1:10){ # Loop for folds
  best.fit = regsubsets(Salary~., data=Hitters[folds!=k,],nvmax=19,method="forward")
  for(i in 1:19){ # Loop for sizes of picked feature best subsets
```

```

    pred = predict(best.fit,Hitters[folds==k,], id=i)
    cv.errors[k,i] = mean((Hitters$Salary[folds==k]-pred)^2)
  }
}
rmse.cv = sqrt(apply(cv.errors,2,mean))
plot(rmse.cv, pch=19, type="b")

```



## Ridge regression and the Lasso

We will use the `glmnet` package, which does not use the formula language, so we will set up an `x` and `y`.

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-2
```

```

x=model.matrix(Salary~.-1, data=Hitters) # -1 means no intercept
y=Hitters$Salary

```

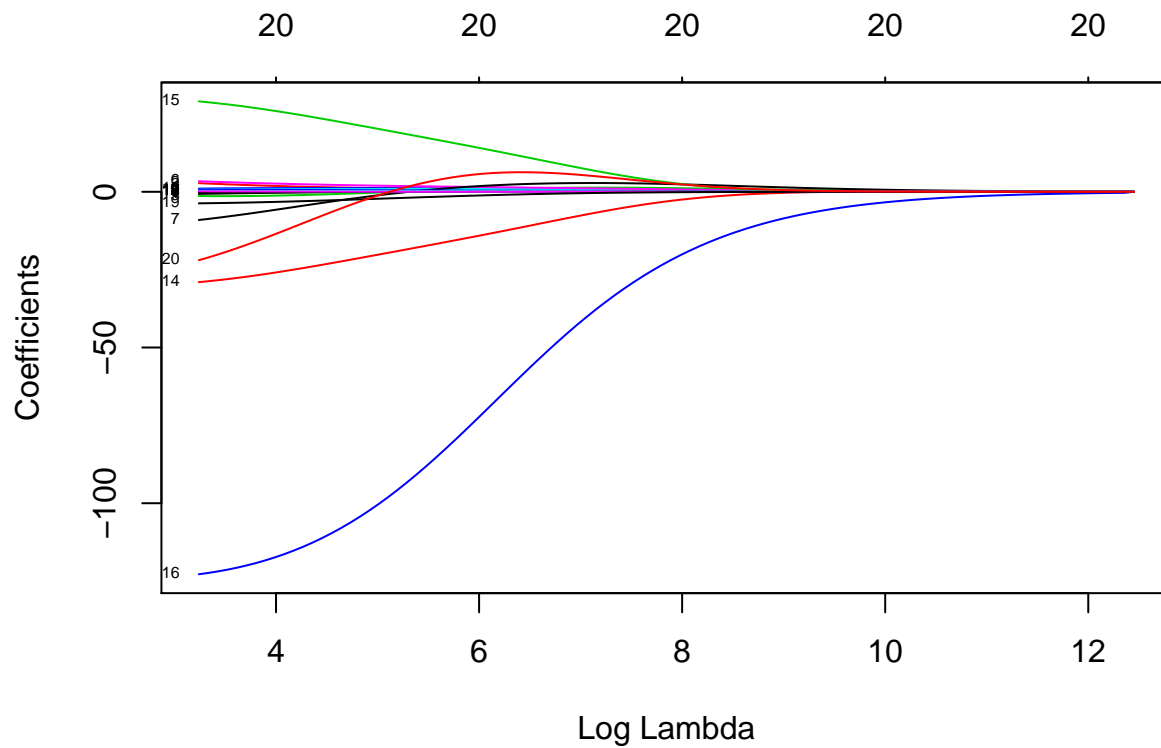
First, we will fit a ridge regression model. This is achieved by calling `glmnet` with `alpha=0` (see the help file). There is also a `cv.glmnet` function, which will do the cross validation for us.

```

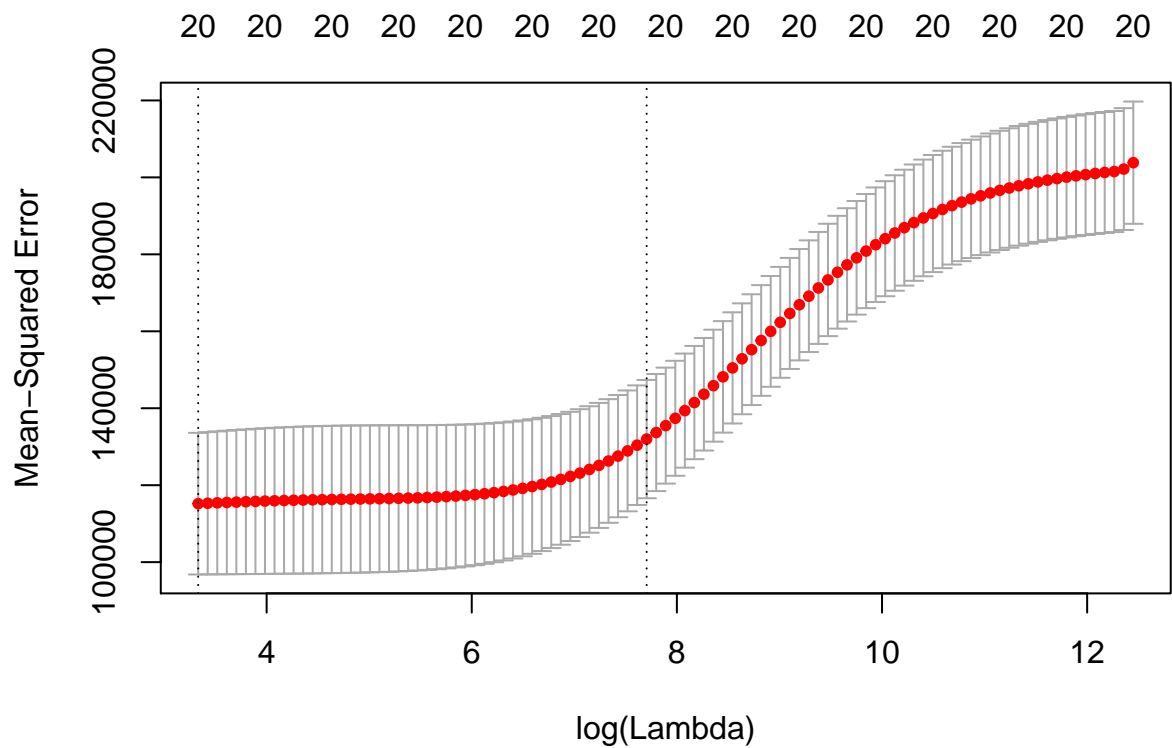
fit.ridge = glmnet(x,y,alpha=0)
plot(fit.ridge, xvar="lambda", label=TRUE)

```



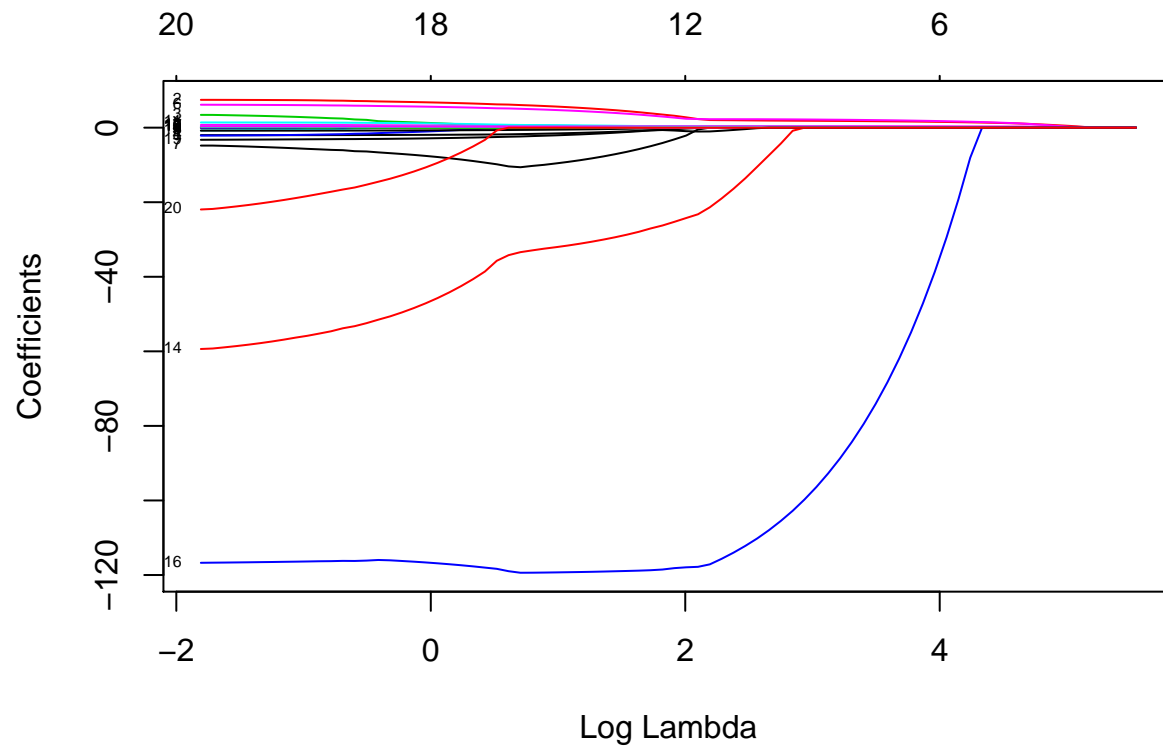


```
cv.ridge = cv.glmnet(x,y,alpha=0)
plot(cv.ridge)
```

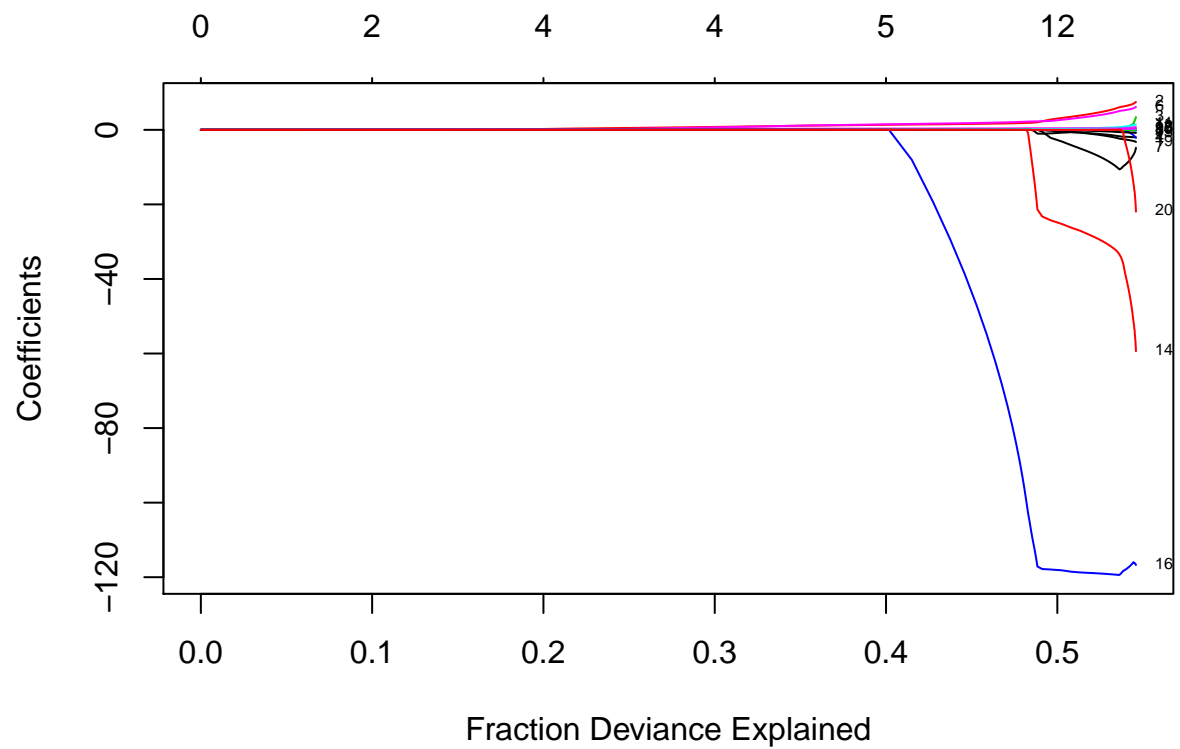


Now, we fit the lasso; in `glmnet` it means `alpha=1`

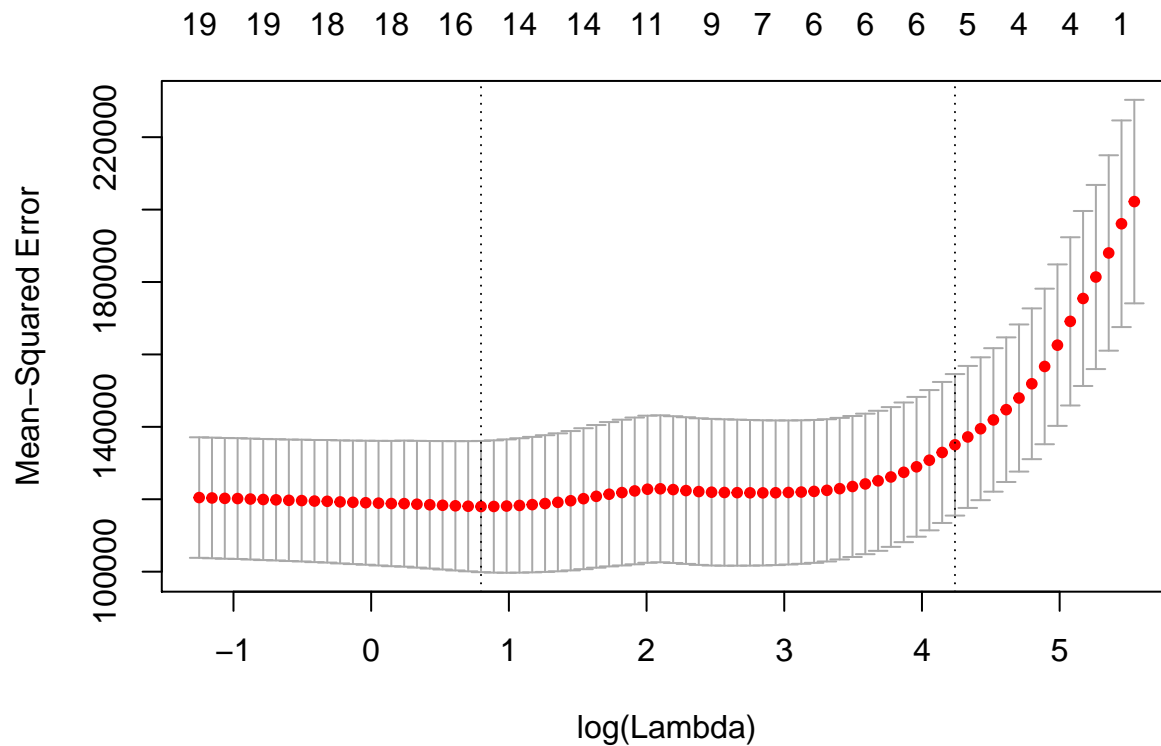
```
fit.lasso = glmnet(x,y)
plot(fit.lasso, xvar="lambda", label=TRUE)
```



```
plot(fit.lasso, xvar="dev", label=TRUE)
```



```
cv.lasso = cv.glmnet(x,y)
plot(cv.lasso)
```



```
coef(cv.lasso) # Picks model one std from minimum to make more parsimonious
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 127.95694754
## AtBat      .
## Hits       1.42342566
## HmRun      .
## Runs       .
## RBI        .
## Walks      1.58214111
## Years      .
## CAtBat     .
## CHits      .
## CHmRun     .
## CRuns      0.16027975
## CRBI       0.33667715
## CWalks     .
## LeagueA    .
## LeagueN    .
## DivisionW  -8.06171262
## PutOuts    0.08393604
## Assists    .
## Errors     .
## NewLeagueN .
```

Suppose we want to use our earlier train/validation division to select the `lambda` for the lasso.

```
lasso.tr = glmnet(x[train,], y[train])  
lasso.tr
```

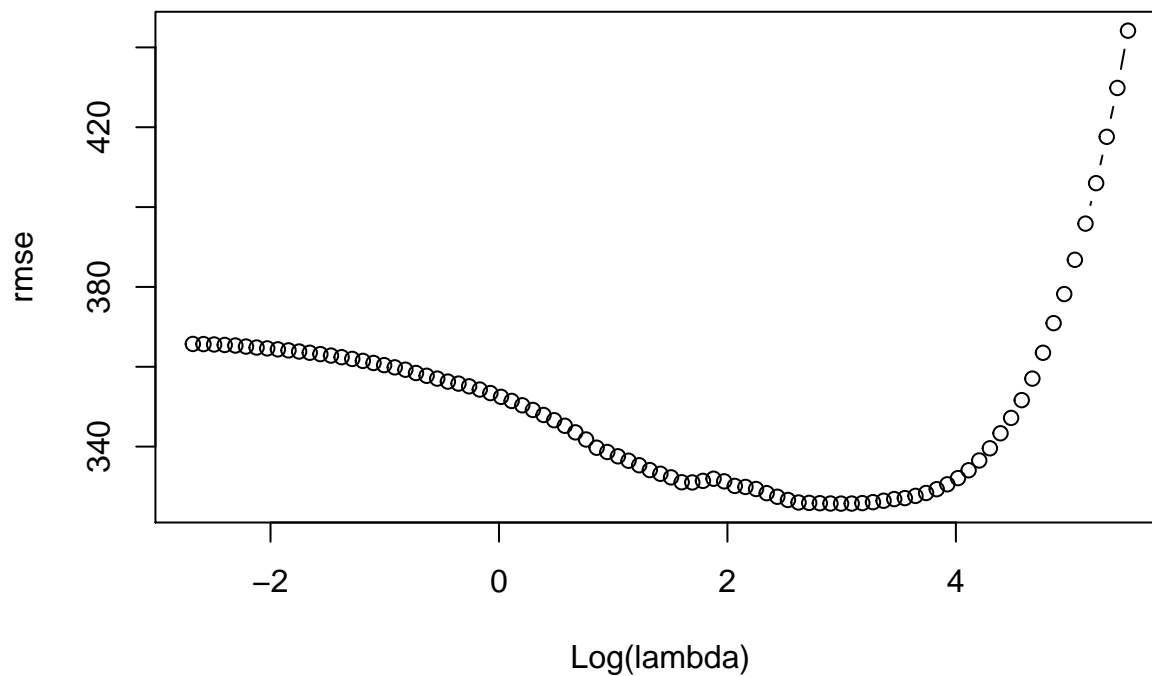
```
##  
## Call:  glmnet(x = x[train, ], y = y[train])  
##  
##           Df      %Dev    Lambda  
## [1,]  0 0.00000 246.40000  
## [2,]  1 0.05013 224.50000  
## [3,]  1 0.09175 204.60000  
## [4,]  2 0.13840 186.40000  
## [5,]  2 0.18000 169.80000  
## [6,]  3 0.21570 154.80000  
## [7,]  3 0.24710 141.00000  
## [8,]  3 0.27320 128.50000  
## [9,]  4 0.30010 117.10000  
## [10,] 4 0.32360 106.70000  
## [11,] 4 0.34310  97.19000  
## [12,] 4 0.35920  88.56000  
## [13,] 5 0.37360  80.69000  
## [14,] 5 0.38900  73.52000  
## [15,] 5 0.40190  66.99000  
## [16,] 5 0.41260  61.04000  
## [17,] 5 0.42140  55.62000  
## [18,] 5 0.42880  50.67000  
## [19,] 5 0.43490  46.17000  
## [20,] 5 0.43990  42.07000  
## [21,] 5 0.44410  38.33000  
## [22,] 5 0.44760  34.93000  
## [23,] 6 0.45140  31.83000  
## [24,] 7 0.45480  29.00000  
## [25,] 7 0.45770  26.42000  
## [26,] 7 0.46010  24.07000  
## [27,] 8 0.46220  21.94000  
## [28,] 8 0.46380  19.99000  
## [29,] 8 0.46520  18.21000  
## [30,] 8 0.46630  16.59000  
## [31,] 8 0.46730  15.12000  
## [32,] 8 0.46810  13.78000  
## [33,] 9 0.47110  12.55000  
## [34,] 9 0.47380  11.44000  
## [35,] 9 0.47620  10.42000  
## [36,] 10 0.48050   9.49500  
## [37,]  9 0.48450   8.65200  
## [38,] 10 0.48770   7.88300  
## [39,] 10 0.49360   7.18300  
## [40,] 11 0.49890   6.54500  
## [41,] 12 0.50450   5.96300  
## [42,] 12 0.51010   5.43400  
## [43,] 13 0.51470   4.95100  
## [44,] 13 0.51850   4.51100  
## [45,] 13 0.52170   4.11000
```

```
## [46,] 14 0.52440 3.74500
## [47,] 14 0.52670 3.41200
## [48,] 15 0.52870 3.10900
## [49,] 15 0.53030 2.83300
## [50,] 15 0.53160 2.58100
## [51,] 16 0.53280 2.35200
## [52,] 17 0.53420 2.14300
## [53,] 18 0.53580 1.95300
## [54,] 18 0.53760 1.77900
## [55,] 18 0.53890 1.62100
## [56,] 18 0.54000 1.47700
## [57,] 18 0.54090 1.34600
## [58,] 18 0.54160 1.22600
## [59,] 18 0.54220 1.11700
## [60,] 18 0.54280 1.01800
## [61,] 18 0.54320 0.92770
## [62,] 18 0.54360 0.84530
## [63,] 18 0.54380 0.77020
## [64,] 19 0.54410 0.70180
## [65,] 19 0.54430 0.63940
## [66,] 19 0.54450 0.58260
## [67,] 19 0.54470 0.53090
## [68,] 19 0.54490 0.48370
## [69,] 20 0.54510 0.44070
## [70,] 20 0.54520 0.40160
## [71,] 20 0.54530 0.36590
## [72,] 20 0.54540 0.33340
## [73,] 20 0.54550 0.30380
## [74,] 20 0.54560 0.27680
## [75,] 20 0.54570 0.25220
## [76,] 20 0.54570 0.22980
## [77,] 20 0.54580 0.20940
## [78,] 20 0.54580 0.19080
## [79,] 20 0.54590 0.17380
## [80,] 20 0.54590 0.15840
## [81,] 20 0.54590 0.14430
## [82,] 20 0.54590 0.13150
## [83,] 20 0.54600 0.11980
## [84,] 19 0.54600 0.10920
## [85,] 19 0.54600 0.09948
## [86,] 19 0.54600 0.09064
## [87,] 19 0.54600 0.08259
## [88,] 20 0.54600 0.07525
## [89,] 20 0.54600 0.06856
```

```
pred = predict(lasso.tr, x[-train,])
dim(pred)
```

```
## [1] 83 89
```

```
rmse = sqrt(apply((y[-train]-pred)^2, 2, mean))
plot(log(lasso.tr$lambda), rmse, type="b", xlab="Log(lambda)")
```



```
lam.best = lasso.tr$lambda[order(rmse)[1]] # Pick best lambda
lam.best
```

```
## [1] 19.98706
```

```
coef(lasso.tr, s=lam.best) # sparse matrix format
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              1
## (Intercept) 107.9416686
## AtBat      .
## Hits       0.1591252
## HmRun      .
## Runs       .
## RBI        1.7340039
## Walks      3.4657091
## Years      .
## CAtBat     .
## CHits      .
## CHmRun     .
## CRuns      0.5386855
## CRBI       .
## CWalks     .
## LeagueA    -30.0493021
## LeagueN    .
## DivisionW  -113.8317016
## PutOuts    0.2915409
## Assists    .
## Errors     .
## NewLeagueN 2.0367518
```